

Customer Management Project

1. Project Overview & Business Goal

This project demonstrates the design and implementation of a modern data pipeline and data warehouse ecosystem.

- ❖ **Objective:** To integrate multiple data sources (customer data, usage logs, transactions, offers, and support tickets), store them in a centralized warehouse, and enable insight generation through PySpark/Hive and business dashboards in Power BI.
- ❖ **Business Goal:** Improve customer churn prediction and lifetime value (CLV) analysis, while providing a single source of truth for reporting.

2. Architecture & Technologies

The architecture follows a data lakehouse pattern, combining raw ingestion, warehouse modeling, and analytics:

- ❖ **Ingestion Layer:** Apache NiFi (streaming ingestion), Apache Sqoop (RDBMS to HDFS)
- ❖ **Storage Layer:** HDFS for raw data, Hive for structured warehouse
- ❖ **Processing Layer:** Hue/Hive for ETL, transformations, and analytics
- ❖ **BI Layer:** Power BI for visualization and insights

3. Data Ingestion Layer

1. **Apache Scoop:** imports relational data (customer) from MariaDB into HDFS and Hive.

```
[student@localhost ~]$ mysql --user=student --password=student
Welcome to the MariaDB monitor.  Commands end with ; or \g.
Your MariaDB connection id is 42
Server version: 5.5.68-MariaDB MariaDB Server

Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

MariaDB [(none)]> create database SIC;
Query OK, 1 row affected (0.00 sec)

MariaDB [(none)]> use SIC;
Database changed
MariaDB [SIC]> CREATE TABLE customers (
  ->   customer_id INT PRIMARY KEY AUTO_INCREMENT,
  ->   first_name VARCHAR(50),
  ->   last_name VARCHAR(50),
  ->   age INT,
  ->   gender VARCHAR(10),
  ->   geography VARCHAR(50),
  ->   is_active TINYINT(1),
  ->   tenure_months INT,
  ->   date_opened DATE
  -> );
Query OK, 0 rows affected (0.01 sec)

MariaDB [SIC]> INSERT INTO customers (first_name, last_name, age, gender, geography, is_active, tenure_mo
nths, date_opened) VALUES
  -> ('Omar','Khalil',34,'Male','Egypt',1,24,'2022-01-15'),
  -> ('Sara','Hassan',29,'Female','UAE',0,12,'2023-01-20'),
  -> ('Ali','Mahmoud',41,'Male','Saudi Arabia',1,36,'2021-05-10'),
  -> ('Noura','Adel',27,'Female','Egypt',1,18,'2022-09-12'),
```

```
[hadoop@localhost ~]$ sqoop import \
> --connect jdbc:mysql://localhost/SIC \
> --username student \
> --password student \
> --table customers \
> --as-parquetfile \
> --target-dir /warehouse/customers
Warning: /usr/local/sqoop/sqoop-1.4.7/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /usr/local/sqoop/sqoop-1.4.7/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
Warning: /usr/local/sqoop/sqoop-1.4.7/./zookeeper does not exist! Accumulo imports will fail.
Please set $ZOOKEEPER_HOME to the root of your Zookeeper installation.
2025-10-02 09:45:10,908 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7
2025-10-02 09:45:10,945 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. C
onsider using -P instead.
2025-10-02 09:45:11,099 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
2025-10-02 09:45:11,099 INFO tool.CodeGenTool: Beginning code generation
2025-10-02 09:45:11,099 INFO tool.CodeGenTool: Will generate java class as codegen_customers
2025-10-02 09:45:11,648 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `customers` AS
t LIMIT 1
2025-10-02 09:45:11,684 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `customers` AS
t LIMIT 1
2025-10-02 09:45:11,690 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/hadoop/hadoop
Note: /tmp/sqoop-hadoop/compile/ce7d979a471e12b7472bd50b1e8fd815/codegen_customers.java uses or overrides
a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
2025-10-02 09:45:14,982 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-hadoop/compile/ce7d979a
471e12b7472bd50b1e8fd815/codegen_customers.jar
2025-10-02 09:45:15,000 WARN manager.MySQLManager: It looks like you are importing from mysql.
2025-10-02 09:45:15,000 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
```

```

025-10-02 09:45:58,972 INFO mapreduce.ImportJobBase: Transferred 21.585 KB in 40.7066 seconds (542.9829
ytes/sec)
025-10-02 09:45:58,975 INFO mapreduce.ImportJobBase: Retrieved 20 records.
hadoop@localhost ~]$ hdfs dfs -ls /warehouse
Found 1 items
-rwxr-xr-x - hadoop supergroup          0 2025-10-02 09:45 /warehouse/customers
hadoop@localhost ~]$ hdfs dfs -ls /warehouse/customers
Found 6 items
-rwxr-xr-x - hadoop supergroup          0 2025-10-02 09:45 /warehouse/customers/.metadata
-rwxr-xr-x - hadoop supergroup          0 2025-10-02 09:45 /warehouse/customers/.signals
-rw-r--r-- 1 hadoop supergroup        2420 2025-10-02 09:45 /warehouse/customers/85b16808-db40-4065-a21a
8b7b4134fe9d.parquet
-rw-r--r-- 1 hadoop supergroup        2432 2025-10-02 09:45 /warehouse/customers/89e74839-1ede-4bf5-a3e7
dffa8587aa19.parquet
-rw-r--r-- 1 hadoop supergroup        2409 2025-10-02 09:45 /warehouse/customers/daa62318-4f2e-4b81-ae30
ebabc4d22ee8.parquet
-rw-r--r-- 1 hadoop supergroup        2434 2025-10-02 09:45 /warehouse/customers/f869c43f-e379-4c04-8c05
8459fde54481.parquet

```

-loading csv files to HDFS:


```


hadoop@localhost ~]$ hdfs dfs -mkdir -p /SIC/tickets_csv
[[A[hadoop@localhost ~]$ hdfs dfs -put /home/hadoop/tickets.csv /SIC/tickets_csv/
hadoop@localhost ~]$ hdfs dfs -put /home/hadoop/tickets.csv /warehouse/tickets_csv/
put: `/warehouse/tickets_csv/': No such file or directory: `hdfs://localhost:9000/warehouse/tickets_csv'
hadoop@localhost ~]$ hdfs dfs -mkdir -p /warehouse/tickets_csv
hadoop@localhost ~]$ hdfs dfs -ls /warehouse
Found 2 items
-rwxr-xr-x - hadoop supergroup          0 2025-10-02 09:45 /warehouse/customers
-rwxr-xr-x - hadoop supergroup          0 2025-10-02 09:51 /warehouse/tickets_csv
hadoop@localhost ~]$ hdfs dfs -put /home/hadoop/tickets.csv /warehouse/tickets_csv/
hadoop@localhost ~]$ nano /home/hadoop/offers.csv
hadoop@localhost ~]$ hdfs dfs -mkdir -p /warehouse/offers_csv
hadoop@localhost ~]$ hdfs dfs -put /home/hadoop/offers.csv /warehouse/offers_csv/

```

2. **Apache NiFi:** ingests semi-structured / unstructured data (JSON) and pushes into HDFS.

-GetFile Processor:





GetFile

GetFile 1.14.0

org.apache.nifi - nifi-standard-nar

In	0 (0 bytes)	5 min
Read/Write	0 bytes / 0 bytes	5 min
Out	0 (0 bytes)	5 min
Tasks/Time	0 / 00:00:00.000	5 min













-Configuration:

Configure Processor


Stopped

SETTINGS | SCHEDULING | PROPERTIES | COMMENTS

Required field +

Property	Value	
Input Directory	 /home/hadoop/usage_json	
File Filter	 [^\\].*	
Path Filter	 No value set	
Batch Size	 10	
Keep Source File	 true	
Recurse Subdirectories	 false	
Polling Interval	 0 sec	
Ignore Hidden Files	 true	
Minimum File Age	 0 sec	
Maximum File Age	 No value set	
Minimum File Size	 0 B	
Maximum File Size	 No value set	

-ConvertRecord Processor:



ConvertRecord ConvertRecord 1.14.0 org.apache.nifi - nifi-standard-nar		
In	0 (0 bytes)	5 min
Read/Write	0 bytes / 0 bytes	5 min
Out	0 (0 bytes)	5 min
Tasks/Time	0 / 00:00:00.000	5 min

-Configuration:

Configure Processor


Stopped

SETTINGS SCHEDULING PROPERTIES COMMENTS

Required field +

Property	Value
Record Reader	JsonTreeReader →
Record Writer	ParquetRecordSetWriter →
Include Zero Record FlowFiles	true

-PutHDFS Processor:



PutHDFS PutHDFS 1.14.0 org.apache.nifi - nifi-hadoop-nar		
In	0 (0 bytes)	5 min
Read/Write	0 bytes / 0 bytes	5 min
Out	0 (0 bytes)	5 min
Tasks/Time	0 / 00:00:00.000	5 min

-Configuration

Configure Processor

Stopped

SETTINGS

SCHEDULING

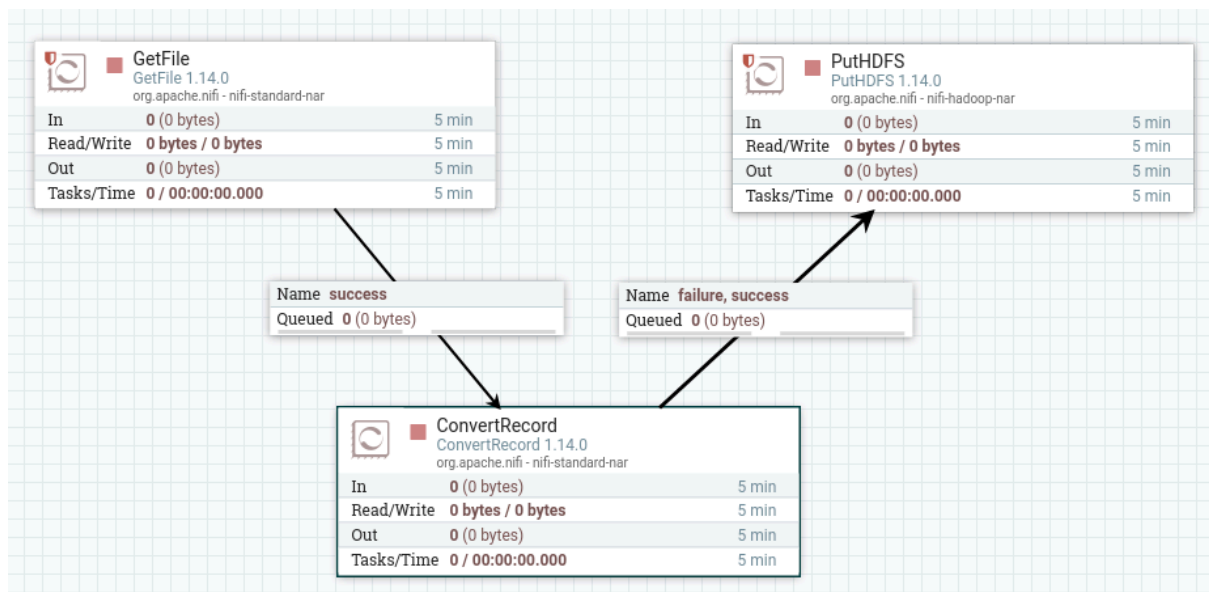
PROPERTIES

COMMENTS

Required field

Property	Value
Hadoop Configuration Resources	No value set
Kerberos Credentials Service	No value set
Kerberos Principal	No value set
Kerberos Keytab	No value set
Kerberos Password	No value set
Kerberos Relogin Period	4 hours
Additional Classpath Resources	No value set
Directory	/hadoop/usage_parquet
Conflict Resolution Strategy	fail

-Full Connection:



4. Mapping External Hive tables to HDFS files:

1. Customers Table:

```
CREATE EXTERNAL TABLE customers (  
    customer_id INT,  
    first_name STRING,  
    last_name STRING,  
    age INT,  
    gender STRING,  
    geography STRING,  
    is_active BOOLEAN,  
    tenure_months INT,  
    date_opened DATE  
)  
STORED AS PARQUET  
LOCATION '/SIC/customers/';
```

2. Customer_support_ticket Table:

```
CREATE EXTERNAL TABLE customer_support_tickets (  
    ticket_id INT,  
    customer_id INT,  
    issue_type STRING,  
    severity STRING,  
    resolution_time_hrs INT  
)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE  
LOCATION '/SIC/customer_support_tickets/';
```

3. Offers Table:

```
CREATE EXTERNAL TABLE offers (  
    offer_id INT,  
    customer_id INT,  
    offer_type STRING,  
    accepted INT,  
    date_offered DATE  
)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE  
LOCATION '/SIC/offers/';
```

4. Usage Table:

```
CREATE EXTERNAL TABLE usage (  
    usage_log_id INT,  
    customer_id INT,  
    product_type STRING,  
    monthly_balance INT,  
    num_products INT  
)  
STORED AS PARQUET  
LOCATION '/SIC/usage/';
```


-Example of how data looks inside hive after mapping tables to HDFS files:

```
INTERMEDIATE:0, RECORDS_OUT_OPERATOR_SEL_1:20,  
2025-10-01 08:51:58,111 INFO [8625e1e5-f110-4edd-9c97-3083952ea55e main] exec.ListSinkOperator: RECORDS_OUT_OPERATOR_LIST_SINK_3:20,  
1001 1 Login Issue Low 5  
1002 2 Card Blocked High 24  
1003 3 Payment Delay Medium 12  
1004 4 Account Locked High 48  
1005 5 App Crash Medium 10  
1006 6 Loan Inquiry Low 6  
1007 7 Transfer Error High 20  
1008 8 ATM Failure Medium 15  
1009 9 Cheque Issue Low 8  
1010 10 KYC Update Medium 14  
1011 11 Card Lost High 30  
1012 12 Loan Delay Medium 18  
1013 13 App Freeze Low 7  
1014 14 Double Deduction High 26  
1015 15 Login Issue Low 6  
1016 16 Card Declined Medium 12  
1017 17 Mobile Token Fail High 28  
1018 18 Password Reset Low 5  
1019 19 App Crash Medium 16  
1020 20 Loan Inquiry Low 9  
Time taken: 0.398 seconds, Fetched: 20 row(s)  
2025-10-01 08:51:58,136 INFO [8625e1e5-f110-4edd-9c97-3083952ea55e main] CliDriver: Time taken: 0.398 seconds, Fetched: 20 row(s)
```

```
INTERMEDIATE:0, RECORDS_OUT_OPERATOR_SEL_1:20,  
2025-10-01 08:43:33,680 INFO [8625e1e5-f110-4edd-9c97-3083952ea55e main] exec.ListSinkOperator: RECORDS_OUT_OPERATOR_LIST_SINK_3:20,  
2001 1 Credit Card 1 2023-02-01  
2002 2 Personal Loan 0 2023-03-12  
2003 3 Savings Account 1 2023-01-10  
2004 4 Cashback 1 2023-05-20  
2005 5 Insurance 0 2023-07-15  
2006 6 Credit Card 1 2023-04-10  
2007 7 Personal Loan 1 2023-02-18  
2008 8 Savings Account 0 2023-06-01  
2009 9 Cashback 1 2023-07-05  
2010 10 Insurance 0 2023-03-21  
2011 11 Credit Card 1 2023-04-12  
2012 12 Personal Loan 1 2023-05-15  
2013 13 Savings Account 0 2023-01-28  
2014 14 Cashback 1 2023-02-22  
2015 15 Insurance 0 2023-06-30  
2016 16 Credit Card 1 2023-05-11  
2017 17 Personal Loan 0 2023-03-05  
2018 18 Savings Account 1 2023-04-19  
2019 19 Cashback 1 2023-07-25  
2020 20 Insurance 0 2023-02-14  
Time taken: 0.276 seconds, Fetched: 20 row(s)  
2025-10-01 08:43:33,726 INFO [8625e1e5-f110-4edd-9c97-3083952ea55e main] CliDriver: Time taken: 0.276 seconds, Fetched: 20 row(s)  
2025-10-01 08:43:33,726 INFO [8625e1e5-f110-4edd-9c97-3083952ea55e main] conf.HiveConf: Using the default value passed in for log id: 8625e1e5-f110-4edd-9c97-3083952ea55e  
2025-10-01 08:43:33,726 INFO [8625e1e5-f110-4edd-9c97-3083952ea55e main] session.SessionState: Resetting thread name to main  
hive> ss
```

4. Data Warehouse and Modeling in Hive

The data warehouse is modeled using a star schema.

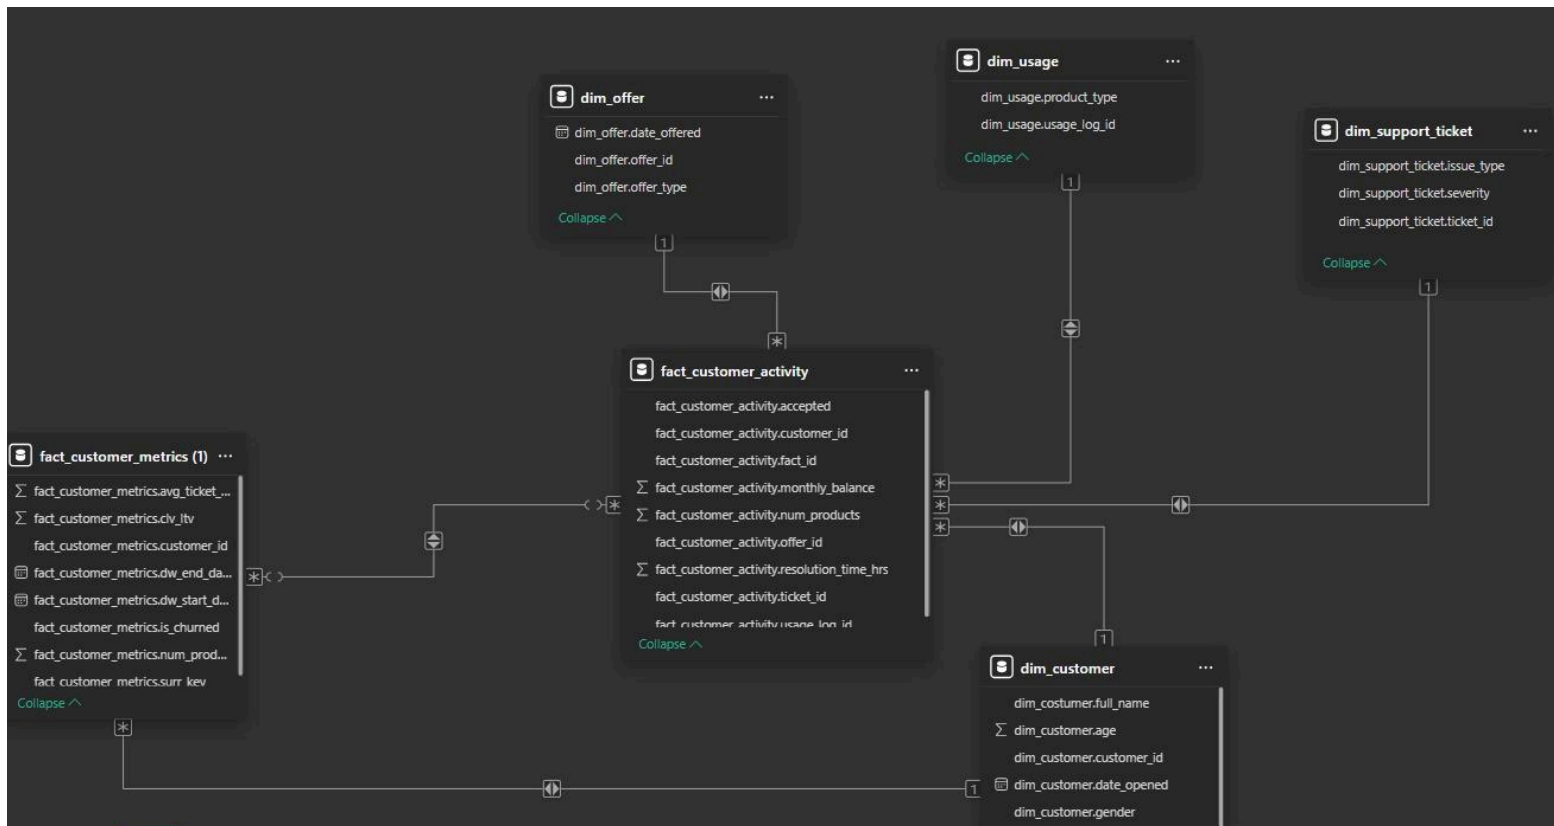
❖ Fact Tables:

- **fact_customer_activity** → contains key customer activity metrics (transactions, usage, churn indicators).
- **fact_customer_metrics_date** → captures customer-level performance metrics such as churn status, customer lifetime value (CLV), average ticket resolution time, number of products, and tracking dates (**dw_start_date**, **dw_end_date**) for historical changes.

❖ Dimension Tables:



- **dim_customer** → customer details, demographics, churn flag.
- **dim_usage** → aggregated usage per period.
- **dim_offer** → marketing campaigns and offers.
- **dim_support_ticket** → customer support interactions.

-Star Schema Diagram:





-Hive Queries:

1. Table dim_customer:

 **Hive**  [Add a name...](#) [Add a description...](#)

```
1 CREATE TABLE dim_customer (  
2     customer_id INT,  
3     first_name STRING,  
4     last_name STRING,  
5     age INT,  
6     gender STRING,  
7     geography STRING,  
8     is_active BOOLEAN,  
9     tenure_months INT,  
10    date_opened DATE  
11 );
```

2. Table dim_offer:

 **Hive**  [Add a name...](#) [Add a description...](#)

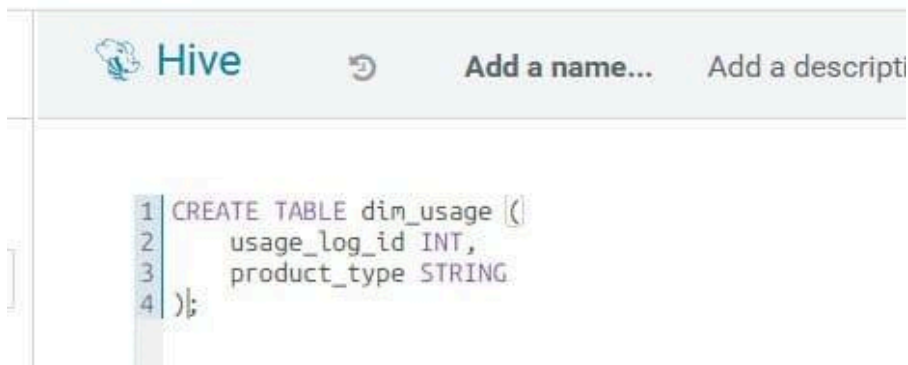
```
1 CREATE TABLE dim_offer (  
2     offer_id INT,  
3     offer_type STRING,  
4     date_offered DATE  
5 );
```

3. Table dim_support_ticket:

 **Hive**  [Add a name...](#) [Add a description...](#)

```
1 CREATE TABLE dim_support_ticket (  
2     ticket_id INT,  
3     issue_type STRING,  
4     severity STRING  
5 );
```

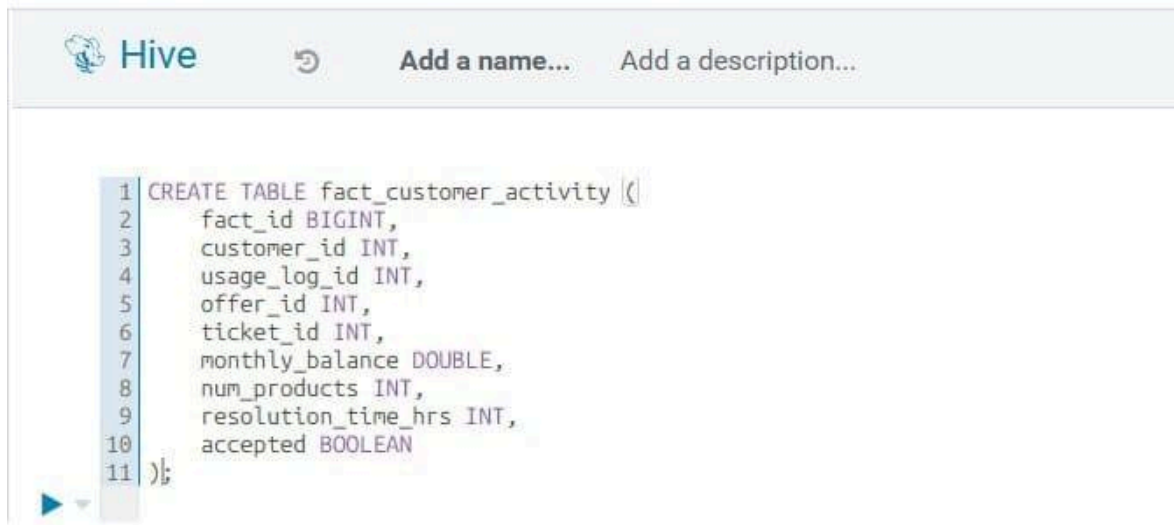
4. Table dim_usage:



The image shows a Hive SQL editor interface. At the top, there is a header bar with the Hive logo, a refresh icon, and two buttons: "Add a name..." and "Add a description...". Below the header, the SQL code for creating the table dim_usage is displayed in a text area with line numbers 1 through 4. The code is as follows:

```
1 CREATE TABLE dim_usage (  
2     usage_log_id INT,  
3     product_type STRING  
4 );
```

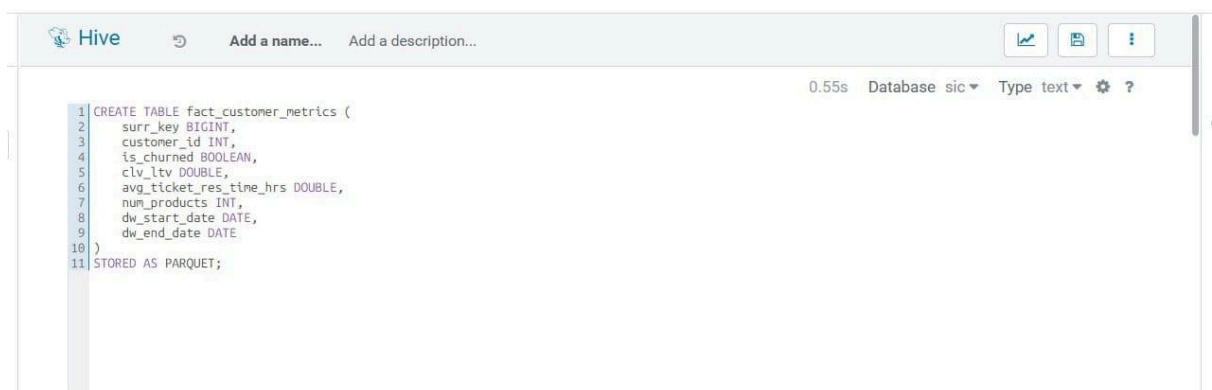
5. Table fact_customer_activity:



The image shows a Hive SQL editor interface. At the top, there is a header bar with the Hive logo, a refresh icon, and two buttons: "Add a name..." and "Add a description...". Below the header, the SQL code for creating the table fact_customer_activity is displayed in a text area with line numbers 1 through 11. The code is as follows:

```
1 CREATE TABLE fact_customer_activity (  
2     fact_id BIGINT,  
3     customer_id INT,  
4     usage_log_id INT,  
5     offer_id INT,  
6     ticket_id INT,  
7     monthly_balance DOUBLE,  
8     num_products INT,  
9     resolution_time_hrs INT,  
10    accepted BOOLEAN  
11 );
```

6. Table fact_customer_metrics:



The image shows a Hive SQL editor interface. At the top, there is a header bar with the Hive logo, a refresh icon, and two buttons: "Add a name..." and "Add a description...". On the right side of the header bar, there are three icons: a line graph, a document, and a help icon. Below the header, the SQL code for creating the table fact_customer_metrics is displayed in a text area with line numbers 1 through 11. The code is as follows:

```
1 CREATE TABLE fact_customer_metrics (  
2     surr_key BIGINT,  
3     customer_id INT,  
4     is_churned BOOLEAN,  
5     clv_ltv DOUBLE,  
6     avg_ticket_res_time_hrs DOUBLE,  
7     num_products INT,  
8     dw_start_date DATE,  
9     dw_end_date DATE  
10 )  
11 STORED AS PARQUET;
```

At the bottom right of the editor, there is a status bar showing "0.55s", "Database sic", "Type text", and icons for settings and help.

-We designed two different fact tables to serve distinct analytical purposes:

- **Fact Customer Activity Table**

This table captures detailed records of individual customer actions. It includes information such as monthly balances, product usage, support tickets, and offers. The purpose of this table is to provide a granular, event-level view of customer interactions, similar to a detailed log or diary of all activities.

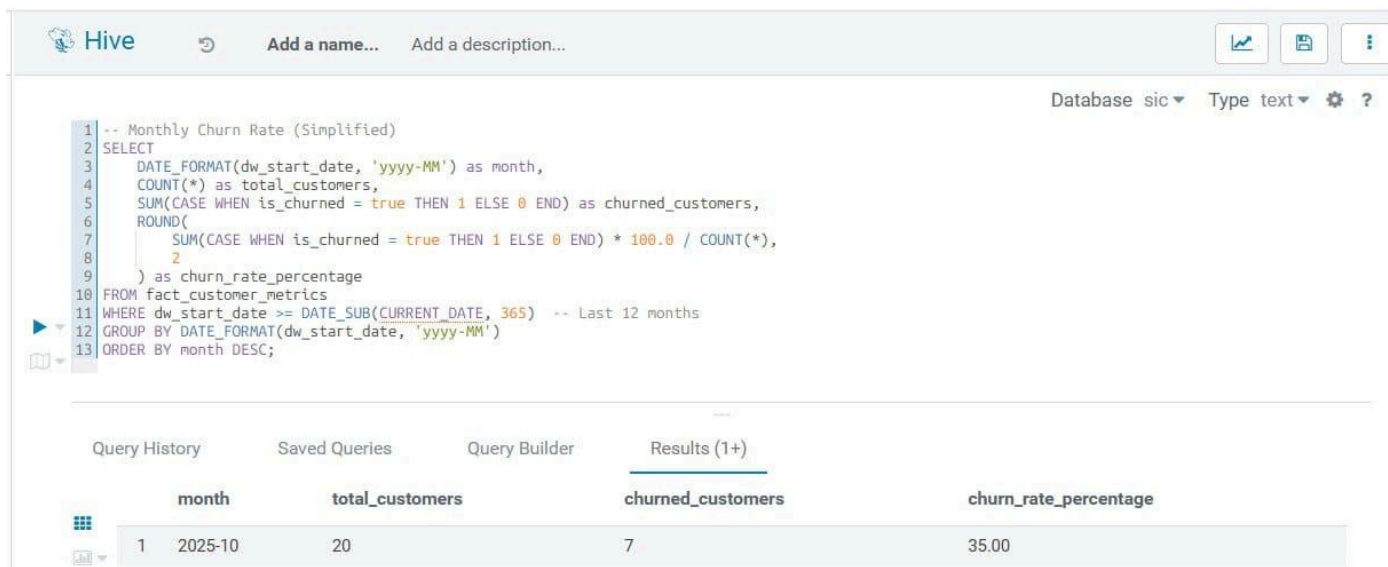
- **Fact Customer Metrics Table**

This table focuses on customer-level summaries and insights derived from the raw activities. It contains calculated metrics such as churn status, customer lifetime value (CLV), and average ticket resolution time. Its role is to provide a consolidated, high-level view of customer performance and risk, similar to a report card.

5. Analysis with Hive (Hue)

Originally, the plan was to use PySpark for advanced analysis, but due to setup issues we performed all the analysis directly in Hive using Hue.

1. Overall customer churn rate for the latest month:



The screenshot shows the Hue web interface for Hive. At the top, there's a header with the 'Hive' logo, a refresh button, and fields for 'Add a name...' and 'Add a description...'. On the right, there are icons for chart, save, and help. Below the header, the query editor shows a SQL query for calculating the monthly churn rate. The query is as follows:

```
1 -- Monthly Churn Rate (Simplified)
2 SELECT
3     DATE_FORMAT(dw_start_date, 'yyyy-MM') as month,
4     COUNT(*) as total_customers,
5     SUM(CASE WHEN is_churned = true THEN 1 ELSE 0 END) as churned_customers,
6     ROUND(
7         SUM(CASE WHEN is_churned = true THEN 1 ELSE 0 END) * 100.0 / COUNT(*),
8         2
9     ) as churn_rate_percentage
10 FROM fact_customer_metrics
11 WHERE dw_start_date >= DATE_SUB(CURRENT_DATE, 365) -- Last 12 months
12 GROUP BY DATE_FORMAT(dw_start_date, 'yyyy-MM')
13 ORDER BY month DESC;
```

Below the query editor, there's a tabbed interface with 'Query History', 'Saved Queries', 'Query Builder', and 'Results (1+)'. The 'Results (1+)' tab is active, displaying a table with the following data:

	month	total_customers	churned_customers	churn_rate_percentage
1	2025-10	20	7	35.00

- **Key Insight:** Our churn rate is critically high. At 35%, we are losing more than one-third of our customers each month. This is an unsustainable business emergency that requires immediate, company-wide action.

2.Churn rate by geography and age group:

```
1 -- Churn Rate by Geography and Age Group
2 SELECT
3     DATE_FORMAT(m.dw_start_date, 'yyyy-MM') as month,
4     c.geography,
5     CASE
6         WHEN c.age < 30 THEN '18-29'
7         WHEN c.age BETWEEN 30 AND 45 THEN '30-45'
8         WHEN c.age BETWEEN 46 AND 60 THEN '46-60'
9         ELSE '60+'
10    END as age_group,
11    COUNT(*) as total_customers,
12    SUM(CASE WHEN m.is_churned = true THEN 1 ELSE 0 END) as churned_customers,
13    ROUND(
14        SUM(CASE WHEN m.is_churned = true THEN 1 ELSE 0 END) * 100.0 / COUNT(*),
15        2
16    ) as churn_rate_percentage
17 FROM fact_customer_metrics m
18 JOIN dim_customer c ON m.customer_id = c.customer_id
19 WHERE m.dw_start_date >= DATE_SUB(CURRENT_DATE, 180) -- Last 6 months
20 GROUP BY
21     DATE_FORMAT(m.dw_start_date, 'yyyy-MM'),
22     c.geography,
23     CASE
24         WHEN c.age < 30 THEN '18-29'
25         WHEN c.age BETWEEN 30 AND 45 THEN '30-45'
26         WHEN c.age BETWEEN 46 AND 60 THEN '46-60'
27         ELSE '60+'
28    END
29 ORDER BY month DESC, churn_rate_percentage DESC;
```

	month	c.geography	age_group	total_customers	churned_customers	churn_rate_percentage
1	2025-10	UAE	18-29	1	1	100.00
2	2025-10	Saudi Arabia	18-29	2	2	100.00
3	2025-10	Qatar	46-60	1	1	100.00
4	2025-10	Qatar	30-45	1	1	100.00
5	2025-10	Bahrain	18-29	1	1	100.00
6	2025-10	Egypt	18-29	3	1	33.33
7	2025-10	UAE	30-45	3	0	0.00
8	2025-10	Saudi Arabia	46-60	1	0	0.00

- **Key Insight:** Churn is concentrated in younger customers (18-29) across all geographies. In several regions, 100% of our young customers churned last month, indicating a critical failure to retain the next generation of clients.

3.Support Impact on Customer Value:

Hive

Add a name...

Add a description...

46.13s Database sic ▼ Type text ▼ ⚙ ?

```
-- Support Impact on Customer Value
SELECT
  t.issue_type,
  t.severity,
  AVG(a.resolution_time_hrs) as avg_resolution_time,
  AVG(m.clv_ltv) as avg_customer_value,
  COUNT(DISTINCT a.customer_id) as affected_customers,
  SUM(CASE WHEN m.is_churned = true THEN 1 ELSE 0 END) as churned_customers,
  ROUND(SUM(CASE WHEN m.is_churned = true THEN 1 ELSE 0 END) * 100.0 / COUNT(DISTINCT a.customer_id), 2) as churn_rate
FROM fact_customer_activity a
JOIN dim_support_ticket t ON a.ticket_id = t.ticket_id
JOIN fact_customer_metrics m ON a.customer_id = m.customer_id
GROUP BY t.issue_type, t.severity
ORDER BY churn_rate DESC;
```

WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

Query History

Saved Queries

Query Builder

Results (17)

	t.issue_type	t.severity	avg_resolution_time	avg_customer_value	affected_customers	churned_customers	churn_rate
1	ATM Failure	Medium	15	20000	1	1	100.00
2	Mobile Token Fail	High	28	288000	1	1	100.00
3	Double Deduction	High	26	33000	1	1	100.00
4	Card Lost	High	30	96000	1	1	100.00
5	Card Blocked	High	24	0	1	1	100.00
6	Loan Inquiry	Low	7.5	42000	2	1	50.00
7	App Crash	Medium	13	82000	2	1	50.00
8	Loan Delay	Medium	18	75000	1	0	0.00
9	KYC Update	Medium	14	0	1	0	0.00
10	Cheque Issue	Low	8	244000	1	0	0.00
11	Account Locked	High	48	63000	1	0	0.00
12	Card Declined	Medium	12	42300	1	0	0.00

- **Key Insight:** High-severity technical and financial issues are a primary driver of churn. Problems like "ATM Failure," "Mobile Token Fail," and "Double Deduction" have a 100% churn rate, directly destroying customer value.

4. Balance behaviour by customer segment

Hive

Add a name...

Add a description...

50.80s

Database sic

Type text

```
1 --Balance Behavior by Customer Segment
2 WITH customer_age_groups AS (
3     SELECT
4         customer_id,
5         geography,
6         CASE
7             WHEN age < 30 THEN '18-29'
8             WHEN age BETWEEN 30 AND 45 THEN '30-45'
9             WHEN age BETWEEN 46 AND 60 THEN '46-60'
10            ELSE '60+'
11        END as age_group
12    FROM dlm_customer
13 )
14 SELECT
15     c.geography,
16     c.age_group,
17     AVG(a.monthly_balance) as avg_monthly_balance,
18     SUM(a.monthly_balance) as total_balance,
19     COUNT(DISTINCT a.customer_id) as active_customers,
20     AVG(m.clv_ltv) as avg_clv,
21     SUM(CASE WHEN a.accepted = true THEN 1 ELSE 0 END) as offers_accepted
22 FROM fact_customer_activity a
23 JOIN fact_customer_metrics m ON a.customer_id = m.customer_id
24 JOIN customer_age_groups c ON a.customer_id = c.customer_id
25 WHERE m.is_churned = false
26 GROUP BY c.geography, c.age_group
27 ORDER BY total_balance DESC;
```

Query History

Saved Queries

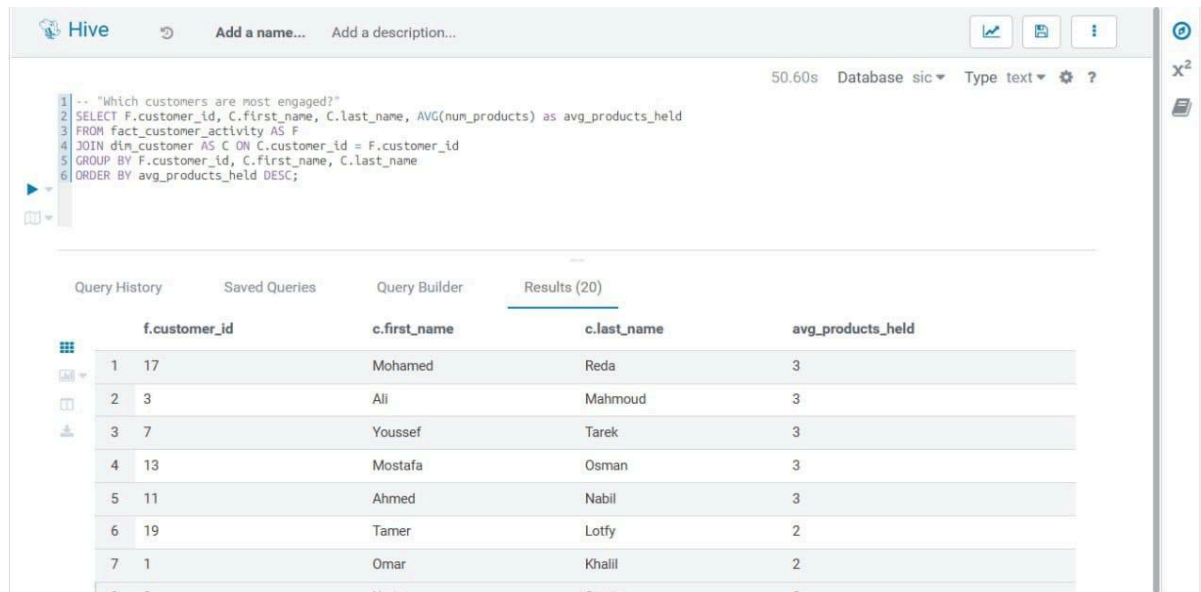
Query Builder

Results (6)

	c.geography	c.age_group	avg_monthly_balance	total_balance	active_customers	avg_clv	offers_accepted
1	UAE	30-45	8666.666666666666	26000	3	295000	2
2	Egypt	30-45	2860	14300	5	65160	4
3	Saudi Arabia	30-45	12000	12000	1	432000	1
4	Egypt	18-29	4100	8200	2	52650	2
5	Kuwait	30-45	6100	6100	1	244000	1
6	Egypt	46-60	0	0	1	0	0

- **Key Insight:** Prime customers are in their 30s-45s living in the UAE and Saudi Arabia. This segment holds the highest average and total balances, has a high CLV, and is most receptive to offers, making them the ideal target for growth initiatives

5. Most engaged customers:

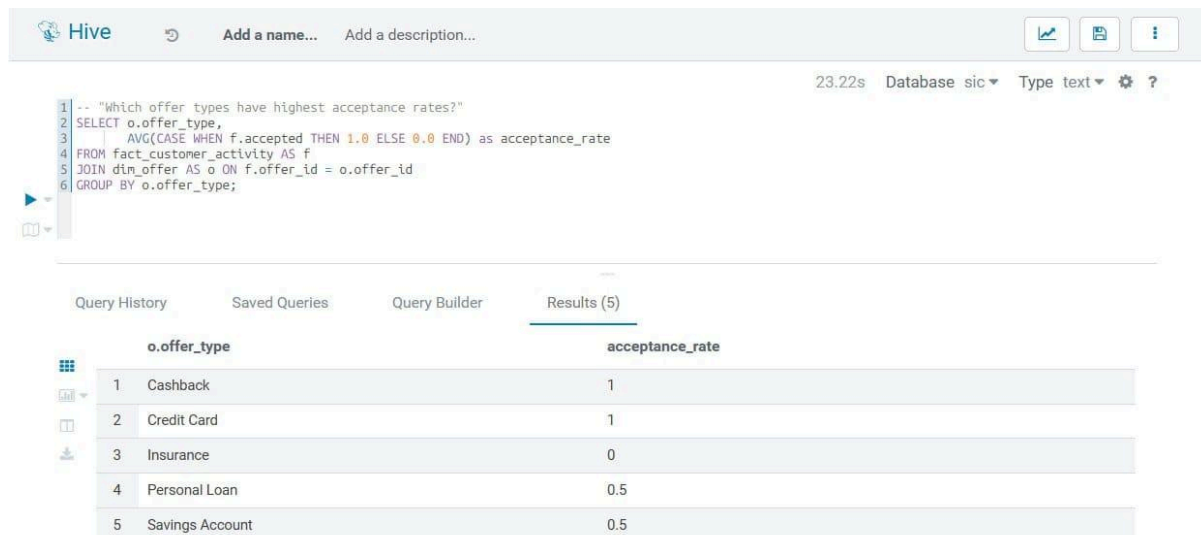


```
1 -- "Which customers are most engaged?"
2 SELECT F.customer_id, C.first_name, C.last_name, AVG(num_products) as avg_products_held
3 FROM fact_customer_activity AS F
4 JOIN dim_customer AS C ON C.customer_id = F.customer_id
5 GROUP BY F.customer_id, C.first_name, C.last_name
6 ORDER BY avg_products_held DESC;
```

	f.customer_id	c.first_name	c.last_name	avg_products_held
1	17	Mohamed	Reda	3
2	3	Ali	Mahmoud	3
3	7	Youssef	Tarek	3
4	13	Mostafa	Osman	3
5	11	Ahmed	Nabil	3
6	19	Tamer	Lotfy	2
7	1	Omar	Khalil	2

- **Key Insight:** These are our most loyal and valuable customers, candidates for loyalty rewards and premium product cross-selling.

6. Offers with highest acceptance rate:

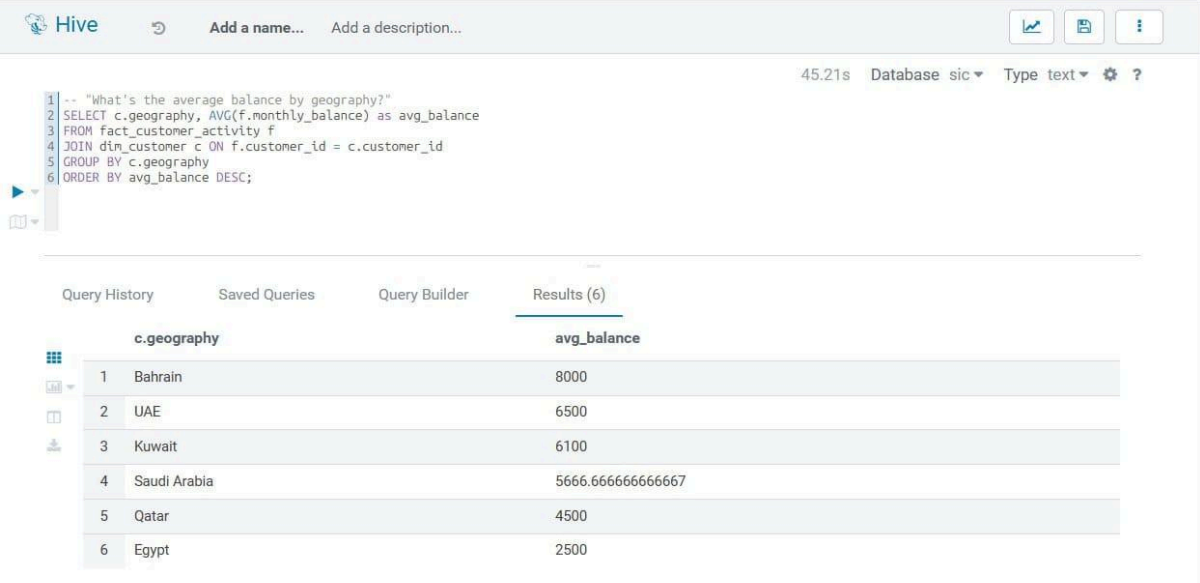


```
1 -- "Which offer types have highest acceptance rates?"
2 SELECT o.offer_type,
3        AVG(CASE WHEN f.accepted THEN 1.0 ELSE 0.0 END) as acceptance_rate
4 FROM fact_customer_activity AS f
5 JOIN dim_offer AS o ON f.offer_id = o.offer_id
6 GROUP BY o.offer_type;
```

	o.offer_type	acceptance_rate
1	Cashback	1
2	Credit Card	1
3	Insurance	0
4	Personal Loan	0.5
5	Savings Account	0.5

- **Key Insight:** Customers prefer immediate, low-risk value. Cashback and Credit Card offers have 100% acceptance, while products like Insurance and Loans see much lower uptake, suggesting customers are wary of long-term commitments.

7. Average balance by geography:



The screenshot shows the Hive query interface. The query is as follows:

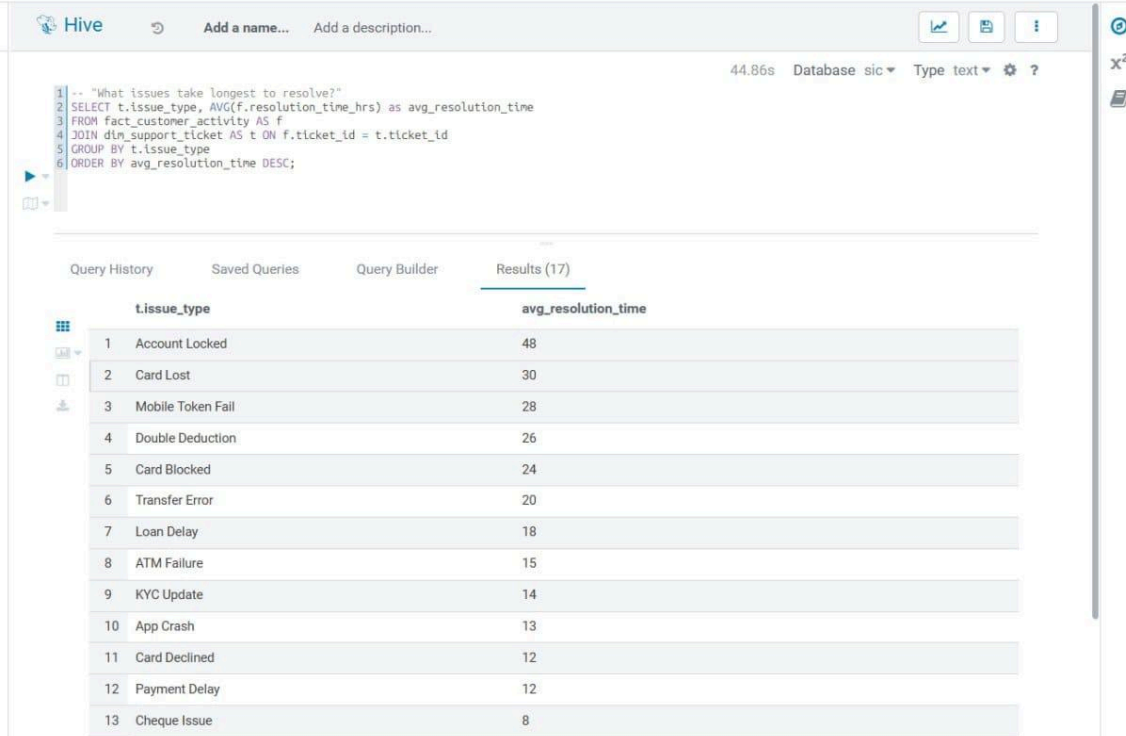
```
1 -- "What's the average balance by geography?"
2 SELECT c.geography, AVG(f.monthly_balance) as avg_balance
3 FROM fact_customer_activity f
4 JOIN dim_customer c ON f.customer_id = c.customer_id
5 GROUP BY c.geography
6 ORDER BY avg_balance DESC;
```

The results are displayed in a table with 6 rows:

c.geography	avg_balance
1 Bahrain	8000
2 UAE	6500
3 Kuwait	6100
4 Saudi Arabia	5666.666666666667
5 Qatar	4500
6 Egypt	2500

- **Key Insight:** Marketing and sales teams can leverage this by tailoring offers to high-balance regions (premium products, cross-selling opportunities), while designing retention-focused initiatives for lower-balance regions (cashback, discounts, loyalty rewards).

8. Issues that take the longest to resolve:



The screenshot shows the Hive query interface. The query is as follows:

```
1 -- "What issues take longest to resolve?"
2 SELECT t.issue_type, AVG(f.resolution_time_hrs) as avg_resolution_time
3 FROM fact_customer_activity AS f
4 JOIN dim_support_ticket AS t ON f.ticket_id = t.ticket_id
5 GROUP BY t.issue_type
6 ORDER BY avg_resolution_time DESC;
```

The results are displayed in a table with 17 rows:

t.issue_type	avg_resolution_time
1 Account Locked	48
2 Card Lost	30
3 Mobile Token Fail	28
4 Double Deduction	26
5 Card Blocked	24
6 Transfer Error	20
7 Loan Delay	18
8 ATM Failure	15
9 KYC Update	14
10 App Crash	13
11 Card Declined	12
12 Payment Delay	12
13 Cheque Issue	8

8.Customer Management Dashboard .

Main Objectives

The main objectives of the dashboard are:

To track customer churn and identify high-risk segments.

To measure Customer Lifetime Value (CLV) and highlight high-value customers.

To analyze product usage patterns across customers.

To evaluate support ticket resolution times by issue type.

To provide actionable insights for improving customer retention and satisfaction.

. Data Model Design

The dashboard uses a star schema consisting of:

Fact Tables:

fact_customer_activity: stores customer activity, usage logs, tickets, and balances.

fact_customer_metrics: contains key metrics such as churn, CLV, number of products.

Dimension Tables:

dim_customer: customer demographics and attributes.

dim_offer: details of offers provided.

dim_usage: product categories used.

dim_support_ticket: issue types and severity.

This design improves query performance and makes analysis flexible.